

Frameworks

Wednesday, October 30, 2024 10:27 AM

Open-Source Frameworks:

AA

Framework	Focus	Key Features	Applications	Strengths	Limitations
Retrieval-Augmented Generation Architectures (RAGAs)	Evaluation by combining LLMs with retrieval systems to enhance factuality and relevance	Retrieves data from external sources during response generation	Applications requiring current or domain-specific knowledge, e.g., customer support, legal, medical	Reduces hallucinations, improves accuracy by grounding in real-time data	Performance depends on retrieval quality and data source accuracy
VADER	Good solution for sentiment analysis of short text	- Rule-based and lexicon approach with valence scores for individual words - Handles degree modifiers, negation, capitalization, and punctuation - Outputs scores for positive, neutral, negative, and compound (overall) sentiment	Can apply to dimensions: toxicity detection, bias detection, tone identification	Performs well on social media posts, informal text	- Performance depends on retrieval quality and data source accuracy - Not as useful on long texts
TruLens	Open-source tool for evaluation, monitoring, and observability	Tracks metrics like relevance, accuracy, fairness, and provides observability features	Continuous monitoring and debugging of production LLMs	Customizable, provides robust feedback loops for model improvement	Requires configuration; complex for large-scale systems
PromptFoo	Evaluation framework for prompt optimization and A/B testing	Allows for prompt experimentation, scoring, and benchmarking across different prompt variations	Fine-tuning and debugging of prompts for specific use cases	Facilitates iterative prompt tuning and side-by-side comparison	Limited to prompt evaluation, not full-model performance
MLflow	Manages model tracking, experimentation, and deployment for ML and LLMs	Tracks experiments, metrics, and artifacts across LLM projects	Large-scale LLM lifecycle management in production	Excellent for reproducibility and lifecycle tracking	General ML tool, not specialized for LLMs
DeepEval	Framework for multi-metric evaluation, including qualitative metrics like interpretability	Supports detailed qualitative analysis of LLMs, including bias, fairness, and interpretability	Applications where interpretability and ethical concerns are high	Comprehensive qualitative metrics for in-depth LLM evaluation	Early-stage, may lack support for complex production workflows
OpenAI Evals	Open-source framework for custom LLM benchmarking	Allows design of custom benchmarks and use-case-specific evaluations	Customizable for specific LLM use cases like support chatbots, QA systems	Extensible, good for specific case-based benchmarking	Still evolving; limited pre-built benchmarks

LCB

Proprietary Frameworks:



Framework	Focus	Key Features	Applications	Strengths	Limitations
ChainPoll	Evaluation of LLMs through stepwise reasoning and consistency polling	Combines chain-of-thought prompting with multiple response polling; averages results to gauge confidence and uncertainty	Evaluating LLM output quality, especially in contexts requiring rationale and consistency checks	Captures nuanced model behavior, provides confidence scores, scales well across contexts	Proprietary, high compute resources for multiple response polling, less effective for simple tasks
Amazon Bedrock	Managed service to evaluate, deploy, and customize multiple foundation models	Access to models from Amazon Titan, Anthropic, AI21 Labs; supports fine-tuning and prompt customization	Multi-model applications in e-commerce, finance, real-time insights	Simplifies customization, deployment, and scaling within AWS	Limited to AWS ecosystem; may lack niche models for specific tasks
NVIDIA NeMo	Evaluation framework for training and fine-tuning LLMs on custom datasets	Offers model training, fine-tuning, and evaluation with synthetic data generation	Custom language tasks, especially in voice and text applications	High compatibility with NVIDIA hardware, supports domain-specific data fine-tuning	Requires NVIDIA infrastructure and expertise in model fine-tuning
Vertex AI Studio	Google's evaluation and fine-tuning platform for LLMs	Offers a model comparison dashboard, tuning capabilities, and built-in metrics	Large-scale LLM deployments with integration into Google Cloud	Simplifies evaluation, tuning, and deployment with Google Cloud resources	Limited to Google Cloud infrastructure
Azure AI Studio	Microsoft's platform for evaluating and deploying LLMs within Azure ecosystem	Offers evaluation tools, model diagnostics, and deployment options	Enterprise-level LLM projects, often within Microsoft ecosystems	Strong integration with Microsoft services, provides deployment tools	Limited to Azure infrastructure, may not support non-Microsoft models
IBM Foundation Model Evaluation (FM-eval)	IBM's framework for evaluation on enterprise metrics like fairness, safety, and interpretability	Evaluates model fairness, interpretability, robustness, and domain-specific knowledge	Enterprises needing compliance-focused or highly interpretable models	Focused on ethical and enterprise metrics, useful for sensitive applications	Primarily suited for IBM environments or IBM-specific models
Arize Phoenix	Primarily evaluates the accuracy and relevance of information retrieval in LLM-generated answers to ensure factual reliability	Customizable templates for assessing response relevance, hallucination detection, and QA correctness in retrieval-heavy tasks	Best suited for applications like customer support, educational tools, and research assistants that depend on accurate information retrieval	Excels in tracking retrieval accuracy, handling complex query responses, and reducing hallucinations, which makes it ideal for QA and RAG applications	Limited flexibility beyond RAG-specific metrics and may require additional (you must pay for) tools for broader LLM evaluation